Rudi Bruchez

From relational to Multimodel Azure Cosmos DB

Yes, it's me

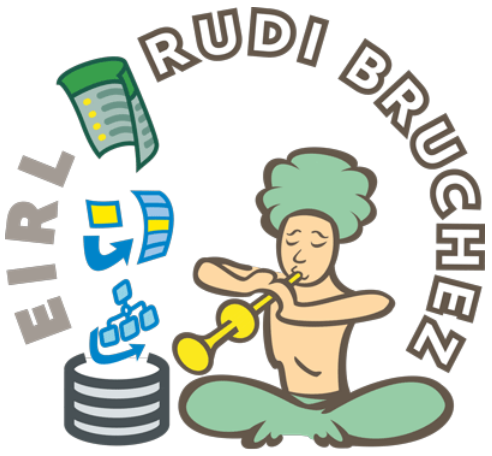

Rudi Bruchez
rudi@babaluga.com
www.babaluga.com

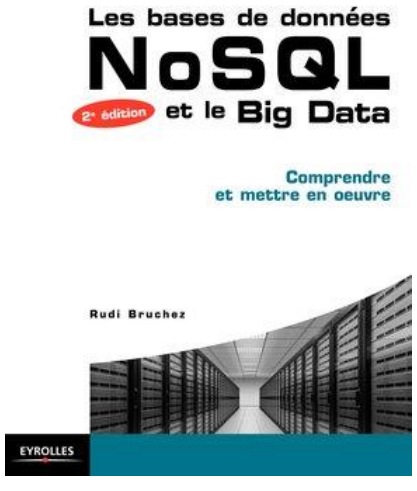

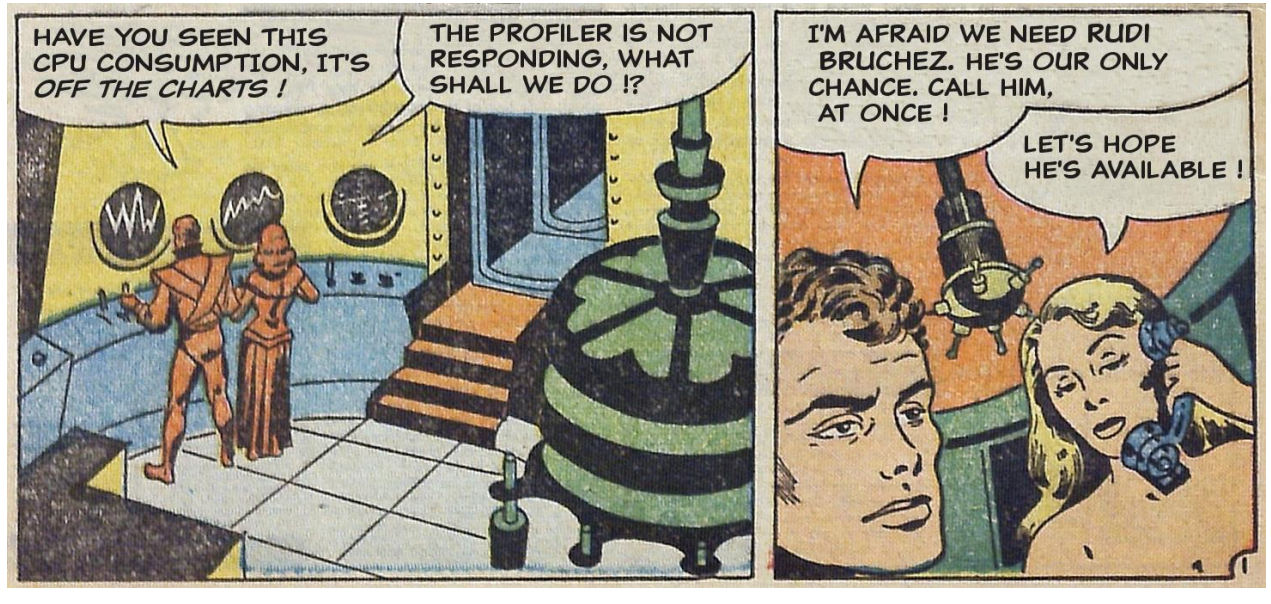


SOLISYON
www.solisyon.de


Zötler Bier
natürlich, sympathisch, frisch


DATAmasterminds


POWER BI SENTINEL
Documentation · Data Lineage · Change Tracking · Backup


ROSEN
empowered by technology


SentryOne


HEDDA.IO


redgate


SHOP APOTHEKE
EUROPE


dbWatch
DATABASE CONTROL

Session Feedback Day 1 (not optional!)

http://bit.ly/DataGrillen2019Day1

Overview

Introduction to CosmosDB

Data Models

Operations

Introduction to CosmosDB

What is Azure CosmosDB

Cosmos DB started in 2010 as “Project Florence”

First availability in 2015 as DocumentDB

Full Cosmos DB released in 2017

What is Azure CosmosDB

Microsoft Azure

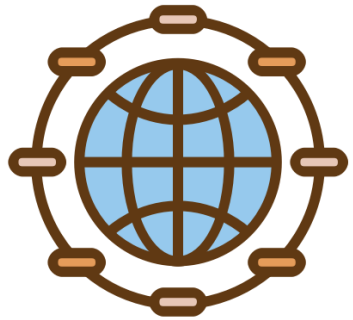

Distributed to regions

Database As A Service

Multi models
Multi APIs

Multi-Model APIs

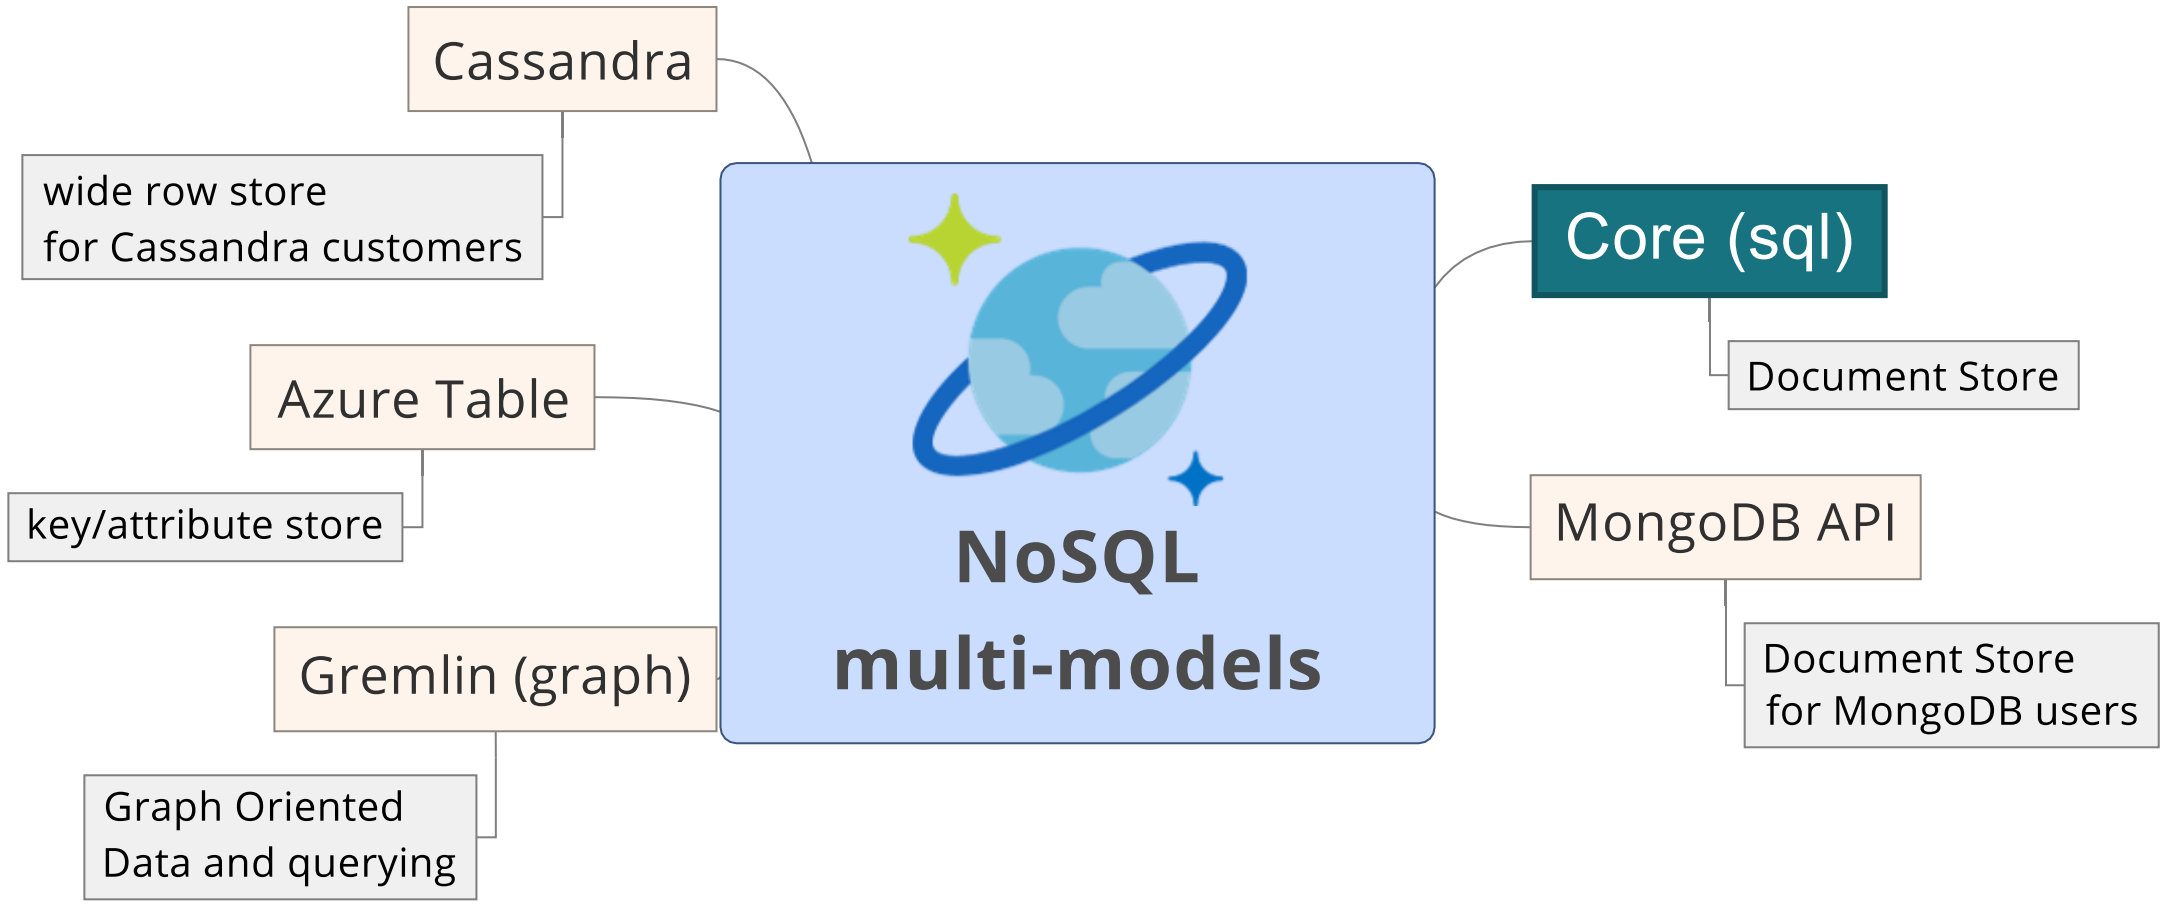

What is the model?

CosmosDB Account

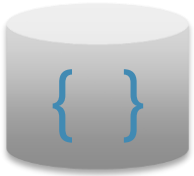

Database

a group of containers

Containers

Contains items

Items

Schema-agnostic data

What is it Good For?

Semi-structured or unstructured data

Document-like data

IOT

Messages processing

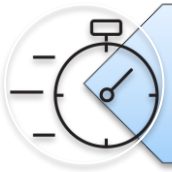

High throuput writes of log or events

Why is it NoSQL?

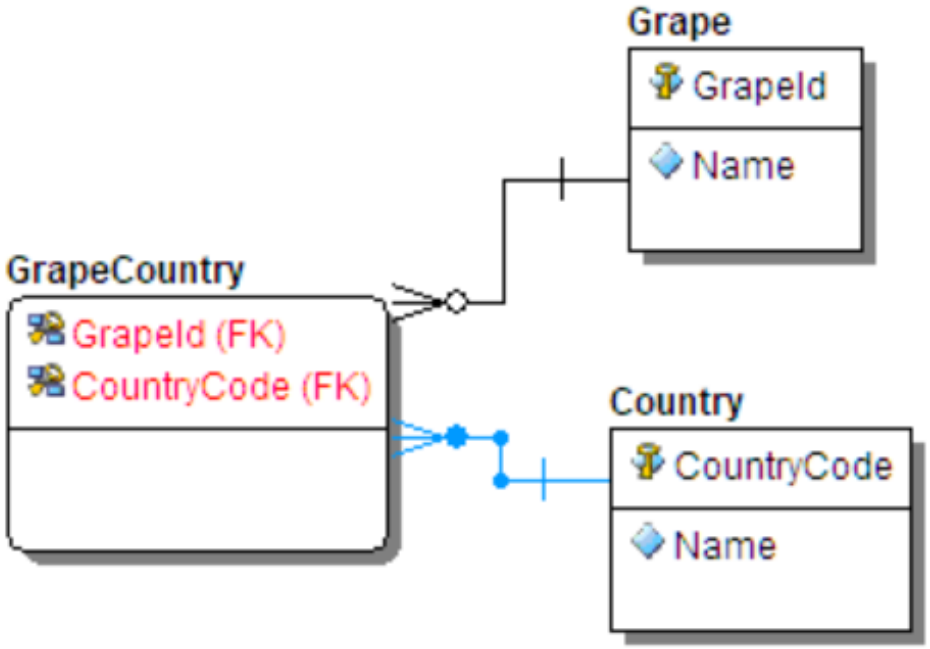

Self-contained items in JSON

document embedding

```
{
    "Grape": "Vionnier",
    "Countries" :
    [
        "France",
        "Germany",
        "Italy"
    ]
}
```

What is in the container?

Containers

Items

Contain fields, internally stored in JSON

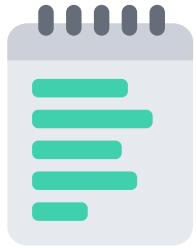

Fields

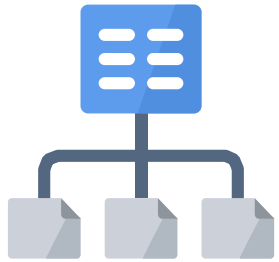

By default, each field is indexed

How you see it in the Data Explorer?

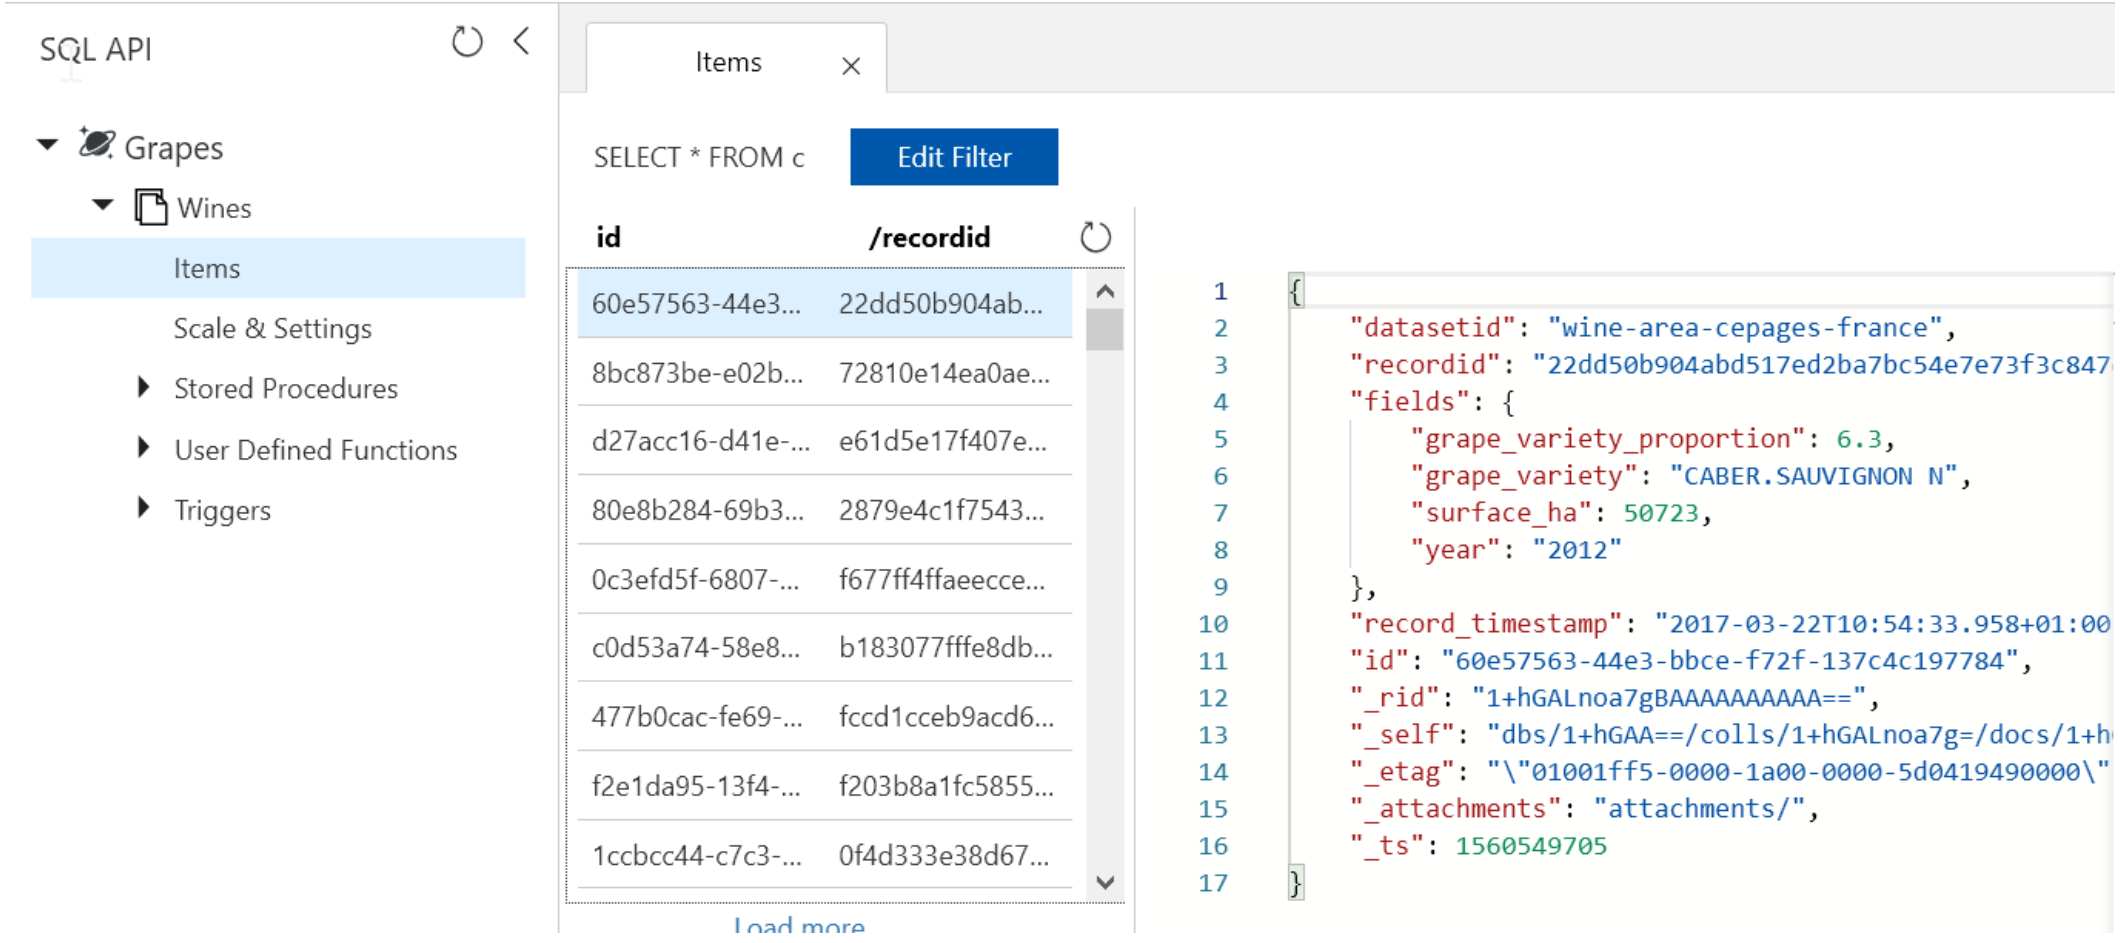

How is it stored?

Schema-agnostic containers

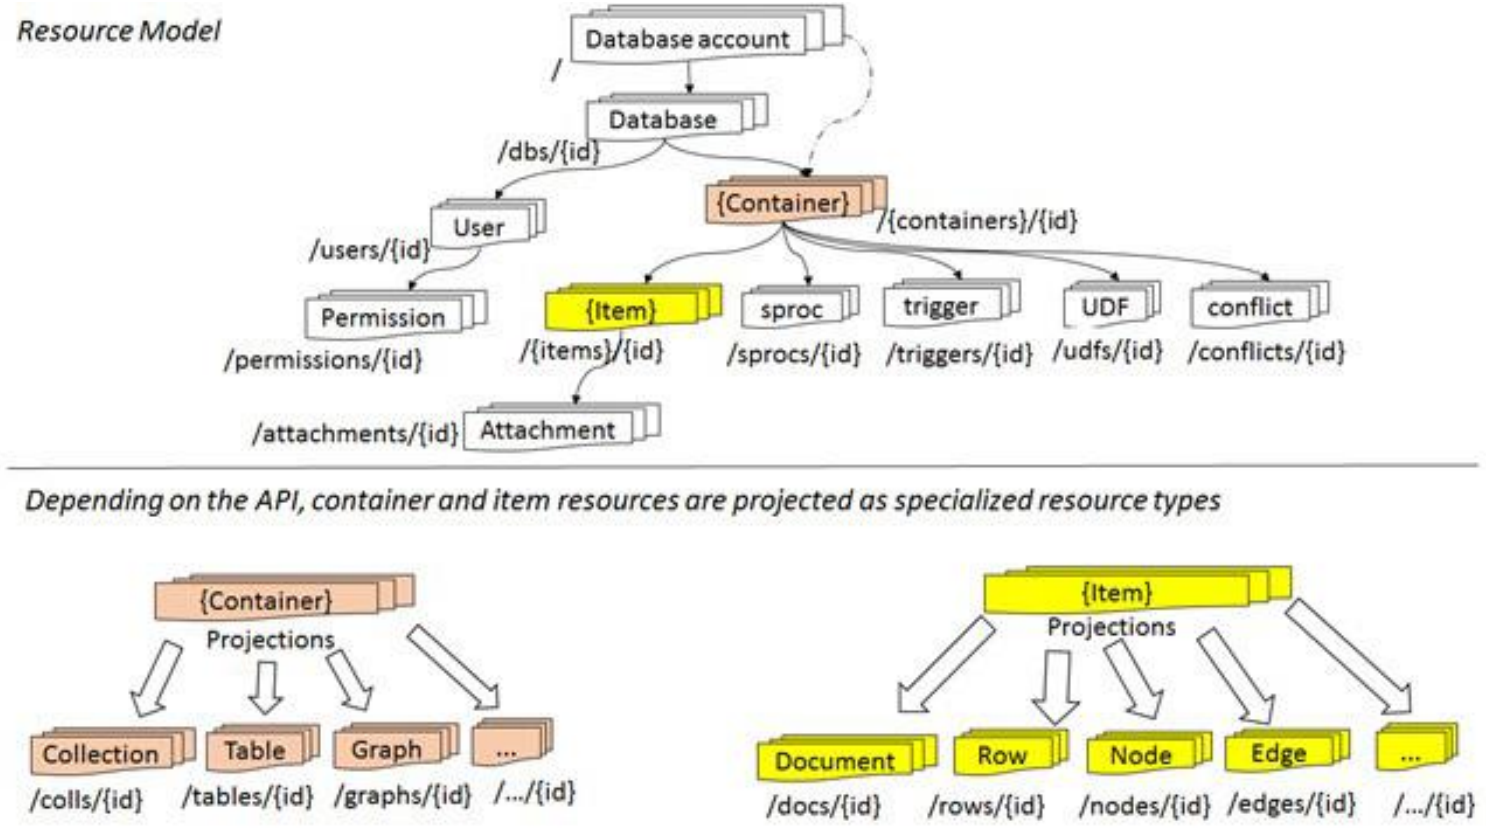

Global distribution

Read-only replicas

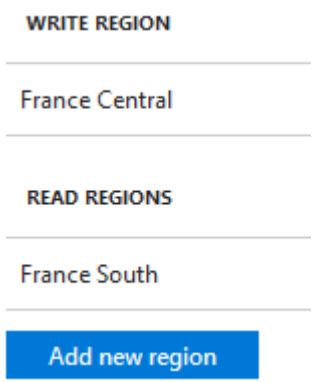

Multi-master for NoSQL API

new Azure Cosmos DB accounts only

September 2018

single digit millisecond write latency at the 99th percentile anywhere in the world

Automatic Partitioning (Sharding)

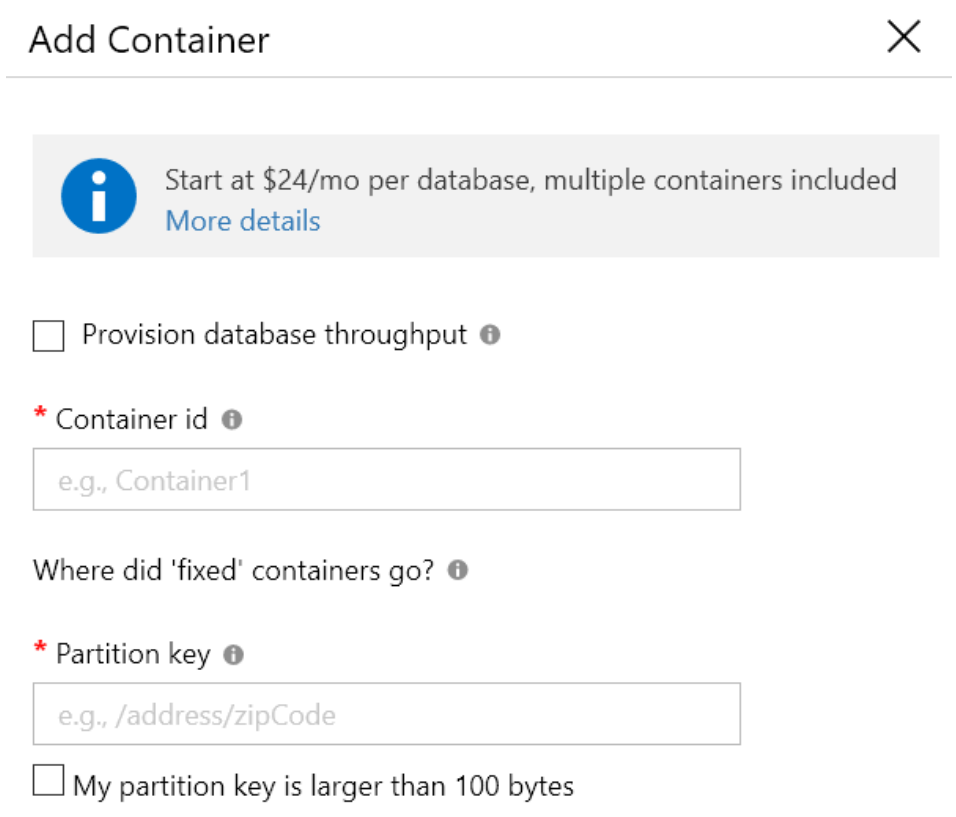

- Logical partitions (vnodes)
- Physical partitioning based on hash is transparent
- Queries in the same partition are more cost-effi
- allowed only in one partition

Model your data from your queries

Partitioning

Automatic physical partitioning if RU/s >= 1000

API	Partition key	Row key
SQL	Custom partition key path	Fixed id
MongoDB	Custom shard key	Fixed _id
Gremlin	Custom partition key property	Fixed id
Table	Fixed PartitionKey	Fixed RowKey

Distributed Consistency?

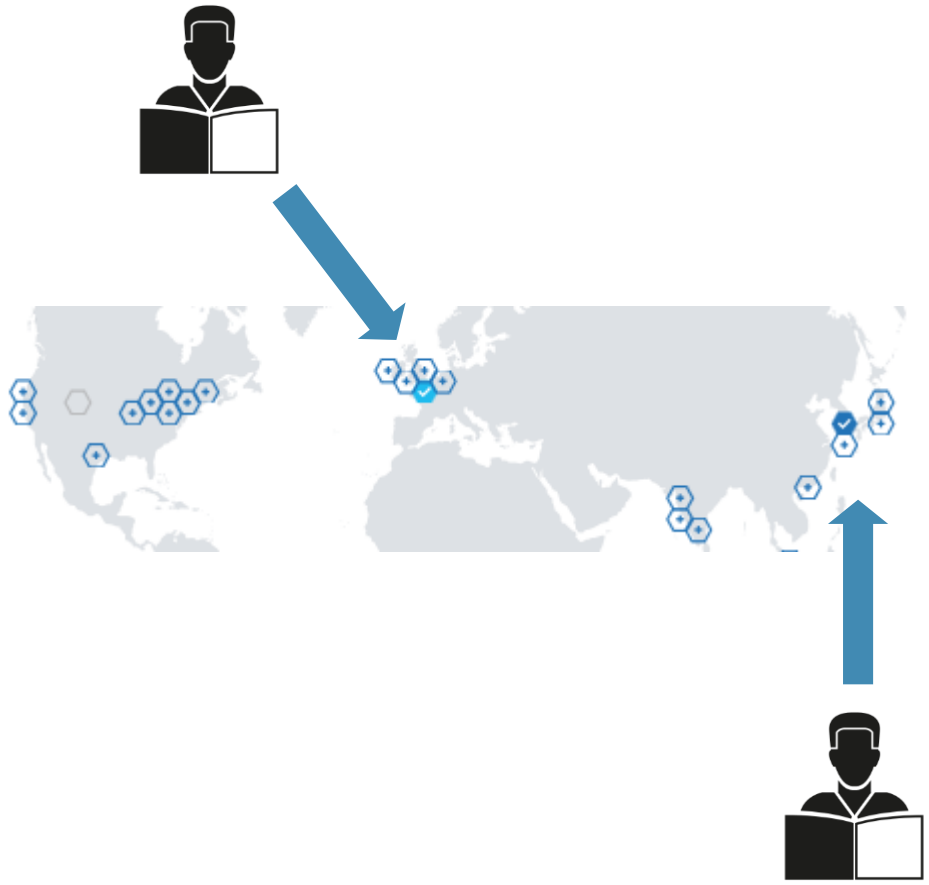

Consistency levels

Consistency level	Description	%
Eventual	No read consistency guaranteed	
Consistent prefix	Eventual, with write-order consistency	
Session	RYOW – consistent inside the same session	73
Bounded Staleness	Set a time or # of operations' lag	20
Strong	Only one region	

2x cheaper

What is a Session ?

```
public async Task GetAsync()
        {
            var response = await this.client.ReadDocumentAsync(...));
            string sessionToken = response.SessionToken;

            RequestOptions options = new RequestOptions();
            options.SessionToken = sessionToken;
            var response2 = await
                client.ReadDocumentAsync(..., options);
        }
```

Per-request consistency level

```
Document doc = client.ReadDocumentAsync(
    documentLink,
    new RequestOptions {
        ConsistencyLevel =
        ConsistencyLevel.Eventual
    }
);
```

2% of Azure Cosmos DB tenants

Multi-Master Databases

September 2018; in Preview before

New accounts only

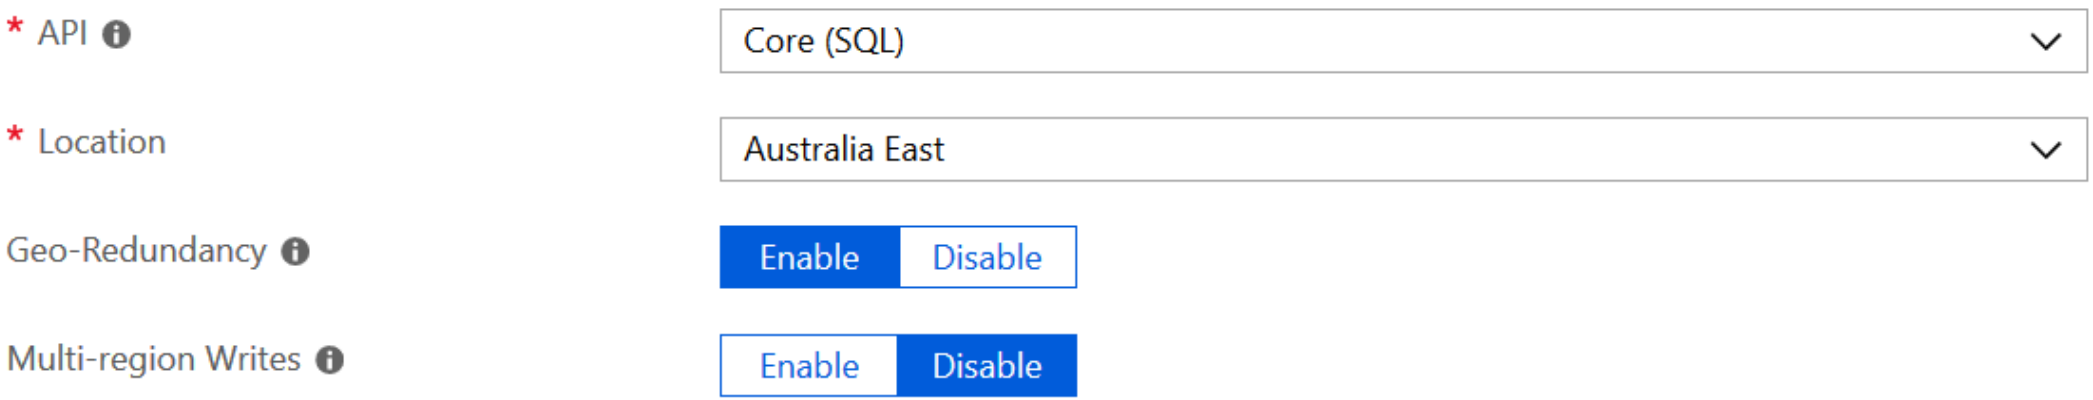

Multi-Master Databases

```
ConnectionPolicy policy = new ConnectionPolicy
{
    ConnectionMode = ConnectionMode.Direct,
    ConnectionProtocol = Protocol.Tcp,
    UseMultipleWriteLocations = true,
};
policy.PreferredLocations.Add("West US");
policy.PreferredLocations.Add("North Europe");
policy.PreferredLocations.Add("Southeast Asia");
```

Multi-homing API

Global conflict resolution modes

Core (SQL) API has 3 modes:

- **Last-Writer-Wins** (LWW) – the largest value in a ConflictResolutionPath wins
- **Custom – User-Defined Procedure** – add an UDP with a special signature to the collection
- **Custom – Asynchronous** – conflicts are not committed, but registered in the read-only conflicts feed for deferred resolution by the application.

For all other API models: LWW only.

Choosing conflict resolution mode

```
DocumentCollection c = await
Client.CreateDocumentCollectionIfNotExistsAsync(
    UriFactory.CreateDatabaseUri("Wines"),
    new DocumentCollection
    {
        Id = "Italy",
        ConflictResolutionPolicy = new ConflictResolutionPolicy
        {
            Mode = ConflictResolutionMode.LastWriterWins,
            ConflictResolutionPath = "/Timestamp"
        }
    }
);
```

https://docs.microsoft.com/fr-fr/azure/cosmos-db/multi-master-conflict-resolution#code-samples

Pricing

By collection

- size
- RU per hour (reserved, not effectively used)

RU – Request Units

- The capacity of your system
- 1 RU = 1 Kb for 1 request

Demos – Creating a Database

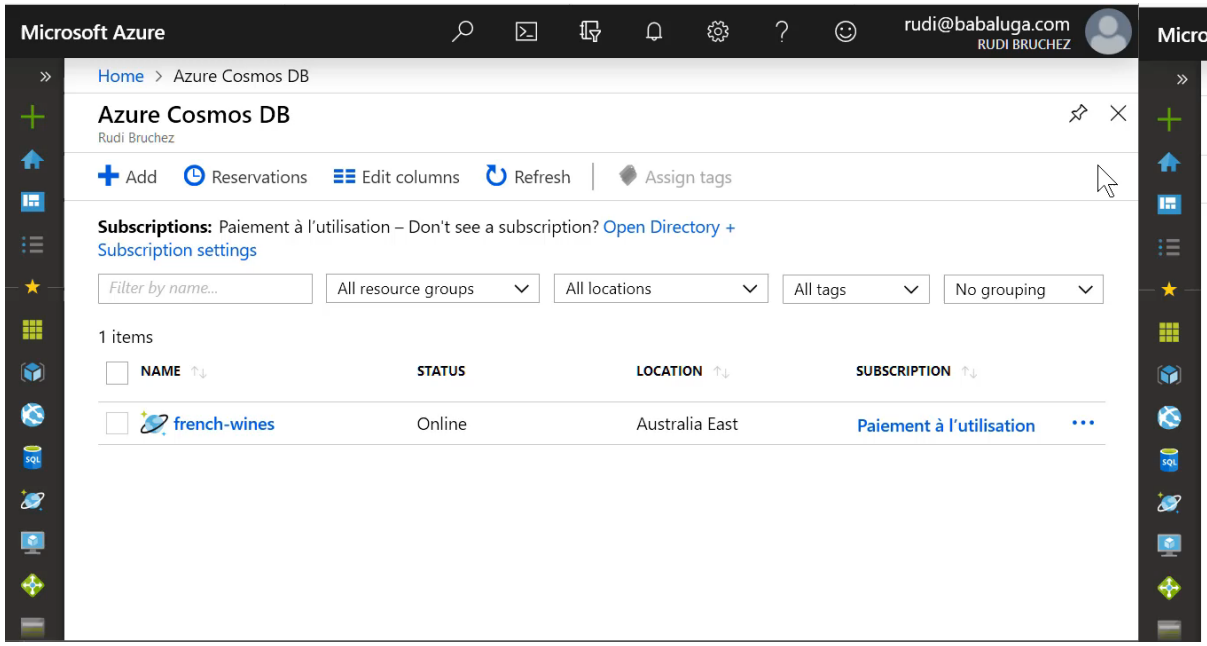

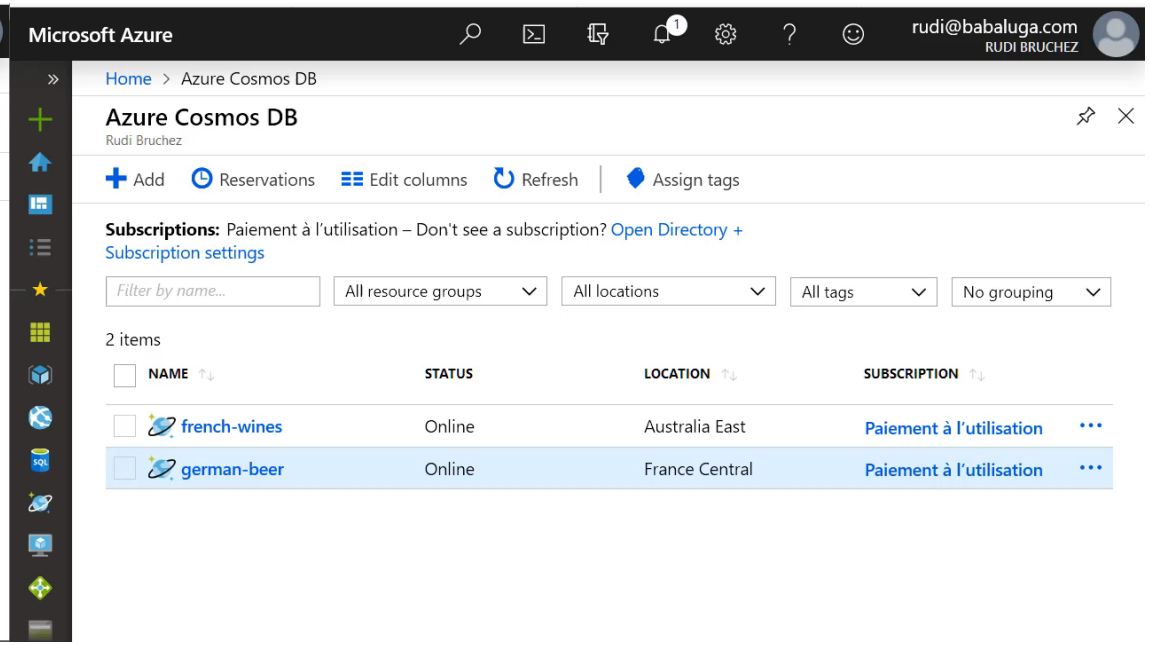

Data Models

Multi-Model APIs

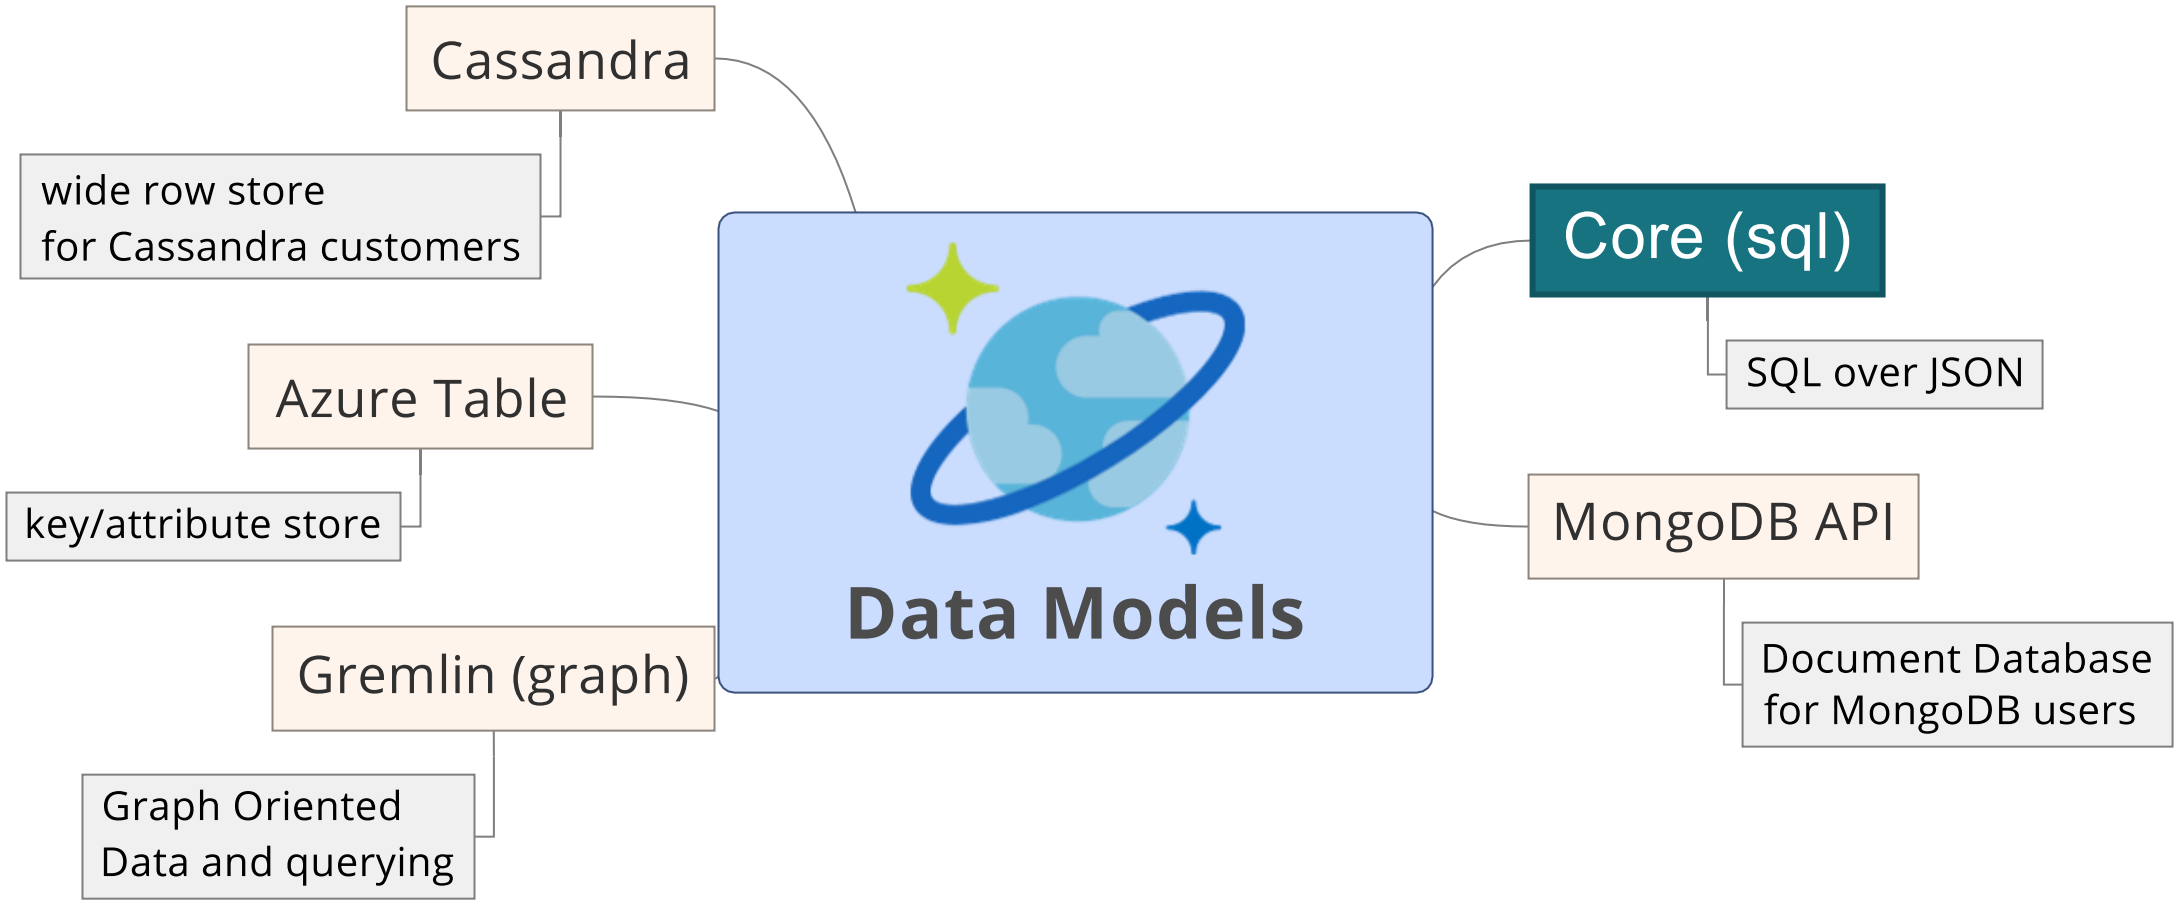

Multi-Models APIs

API	Mapping		Compatibility
Core	Containers	Items	
MongoDB	Collections	Documents	MongoDB 3.2, some 3.4 features in preview
Gremlin	Graphs	Nodes, Edges	Gremlin 3.2
Cassandra	Tables	Rows	CQL 4
Azure Table Storage	Table	Item	

SQL API

Document Database (JSON)

documentDB API

“id” column mandatory, manually or automatically set

SQL language for documents

Demo – Creating an Item

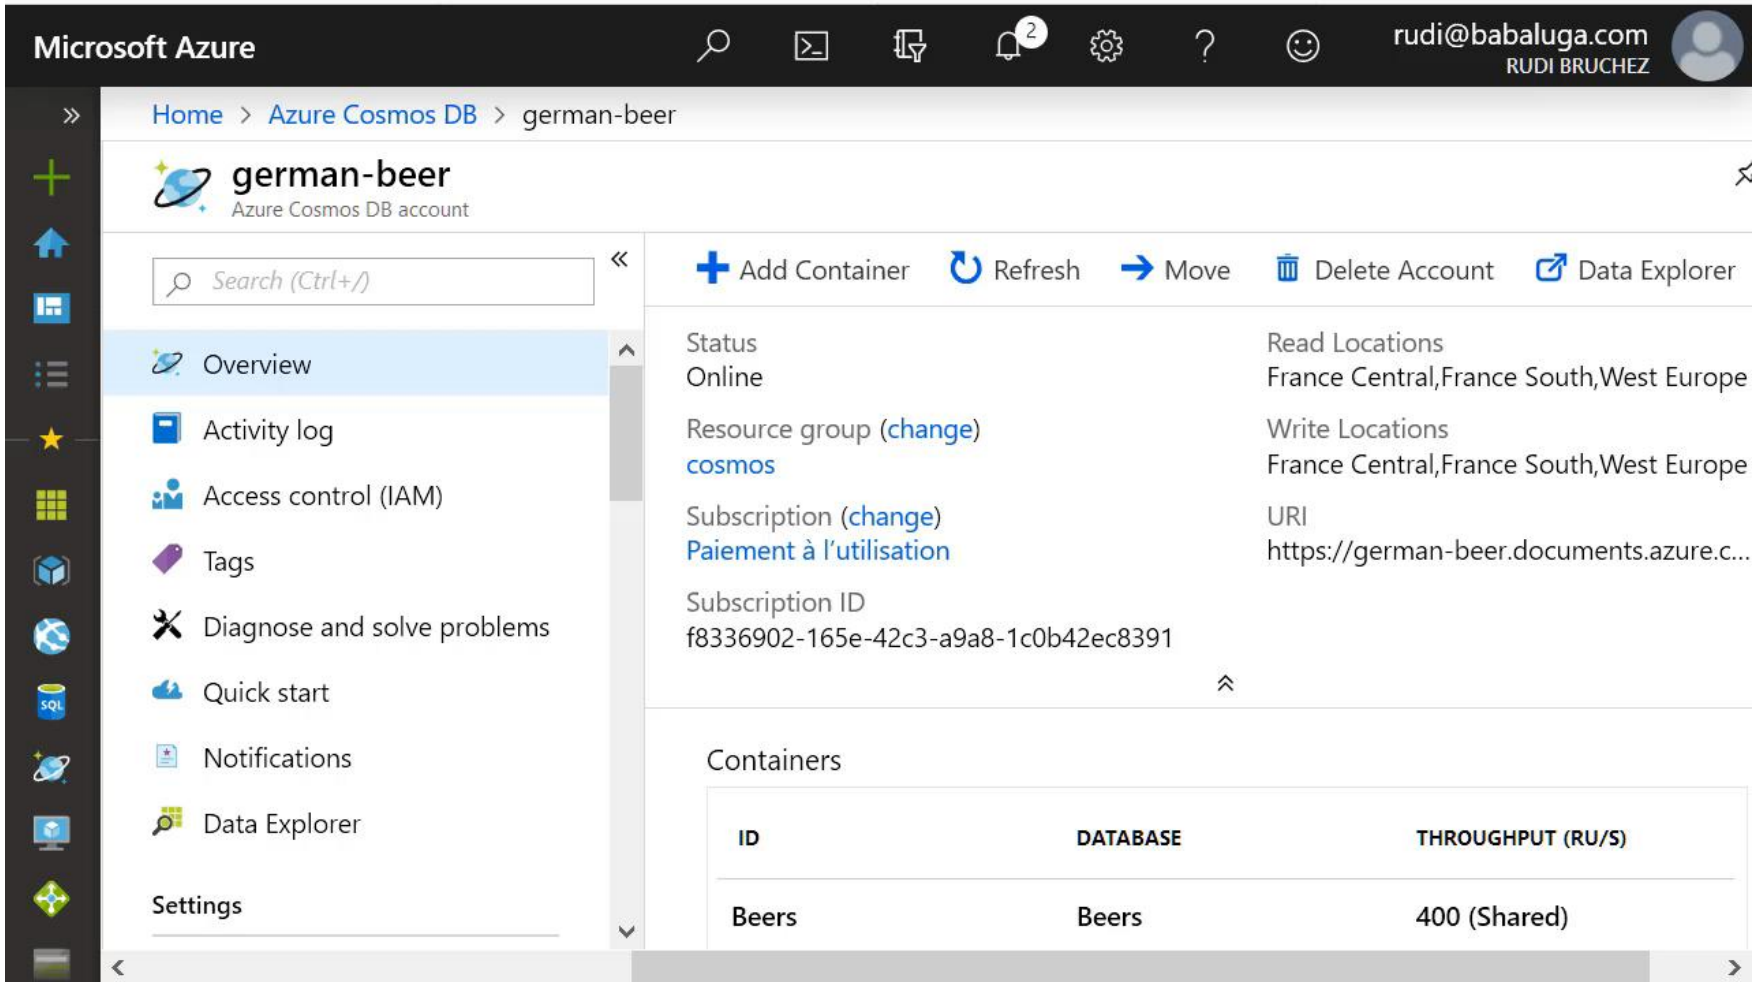

What is in the Document?

Property	User settable ?	Purpose
_rid	System	unique, hierarchical identifier of the resource
_etag	System	for optimistic concurrency control
_ts	System	Last updated timestamp (epoch)
_self	System	Unique addressable URI of the resource
id	Either	If the user does not specify, system generated

eTag Management – OCC, MVCC

```
{
    "id": "AltaMora_EtnaBianco_2017",
    "Name": "Alta Mora, Etna Bianco",
    "Variety": "Carricante",
    "Year": 2017,
    "Country": "Sicily",
    "_rid": "UL0VALLyZwABAAAAAAAAAA==",
    "_self": "dbs/UL0VAA==/colls/UL0VALLyZwA=/do
    "_etag": "\"00003401-0000-0000-0000-5aff136d
    "_attachments": "attachments/",
    "_ts": 1526666093
}
```

```
var ac = new AccessCondition {
    Condition = doc.ETag, Type = AccessConditionType.IfMatch};

this.client.ReplaceDocumentAsync(doc.SelfLink, wine,
    new RequestOptions {AccessCondition = ac}
);
```

060.etag.cs

SQL Queries

- Subset of SQL implemented in Javascript
- Javascript support
- JSON projection
- Intra-document joins
- Support for array iteration in the From clause

SQL Queries

```
SELECT Name as "Wine Name"
FROM "all-wines"
-- returns a JSON list


SELECT {"Wine Name": Name} as "Wines"
FROM "all-wines"
-- returns a JSON object
```

SQL Queries

```
SELECT Name as "Seller Name"
FROM wines.sellers
-- accessing a subdocument


SELECT Name as "Seller Name"
FROM w IN wines.sellers
-- iterating through a JSON array
```

Intra-document joins

- Only INNER JOIN
- Complete cross product of the sets participating in the join

```
SELECT tag.name
FROM food
JOIN tag IN food.tags
WHERE food.id = "09052"
```

Builtin Functions

Function group	Operations
Mathematical functions	ABS, CEILING, EXP, FLOOR, POWER, ROUND, SIGN, SQRT, SQUARE, …
Type checking functions	IS_ARRAY, IS_BOOL, IS_NULL, IS_NUMBER, IS_OBJECT, IS_STRING, IS_DEFINED, and IS_PRIMITIVE
String functions	CONCAT, CONTAINS, ENDSWITH, INDEX_OF, LEFT, LENGTH, LOWER, LTRIM, REPLACE, REPLICATE, REVERSE, RIGHT, RTRIM, STARTSWITH, SUBSTRING, UPPER
Array functions	ARRAY_CONCAT, ARRAY_CONTAINS, ARRAY_LENGTH, and ARRAY_SLICE
Spatial functions	ST_DISTANCE, ST_WITHIN, ST_INTERSECTS, ST_ISVALID, and ST_ISVALIDDETAILED

Demos – SQL API

Importing JSON

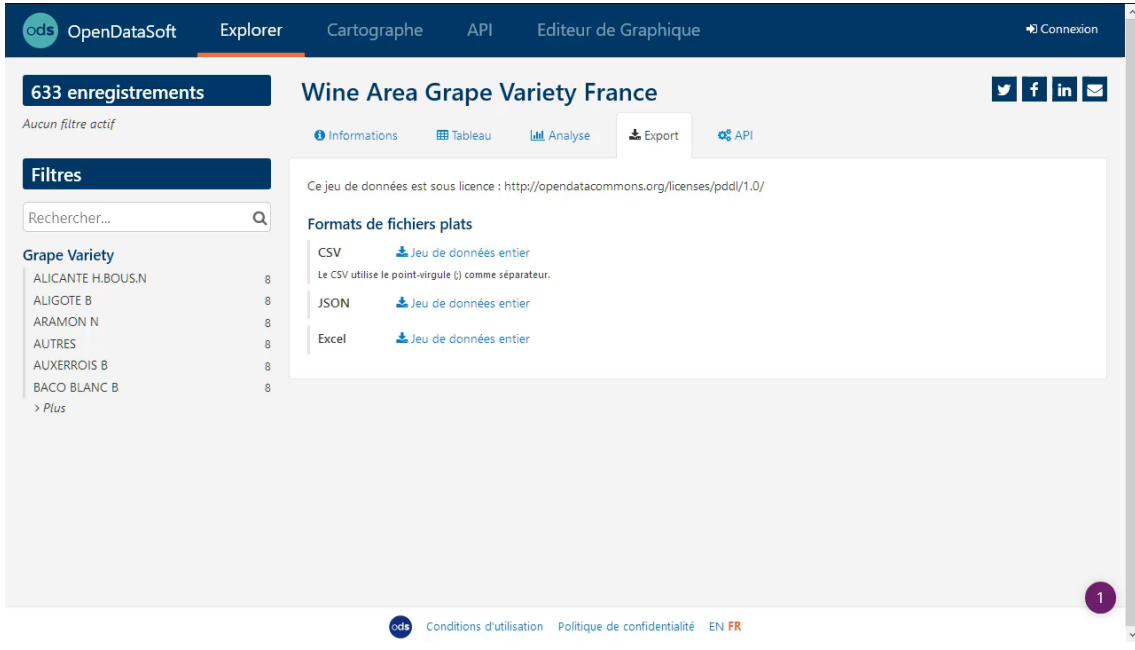

Querying The Wines Container

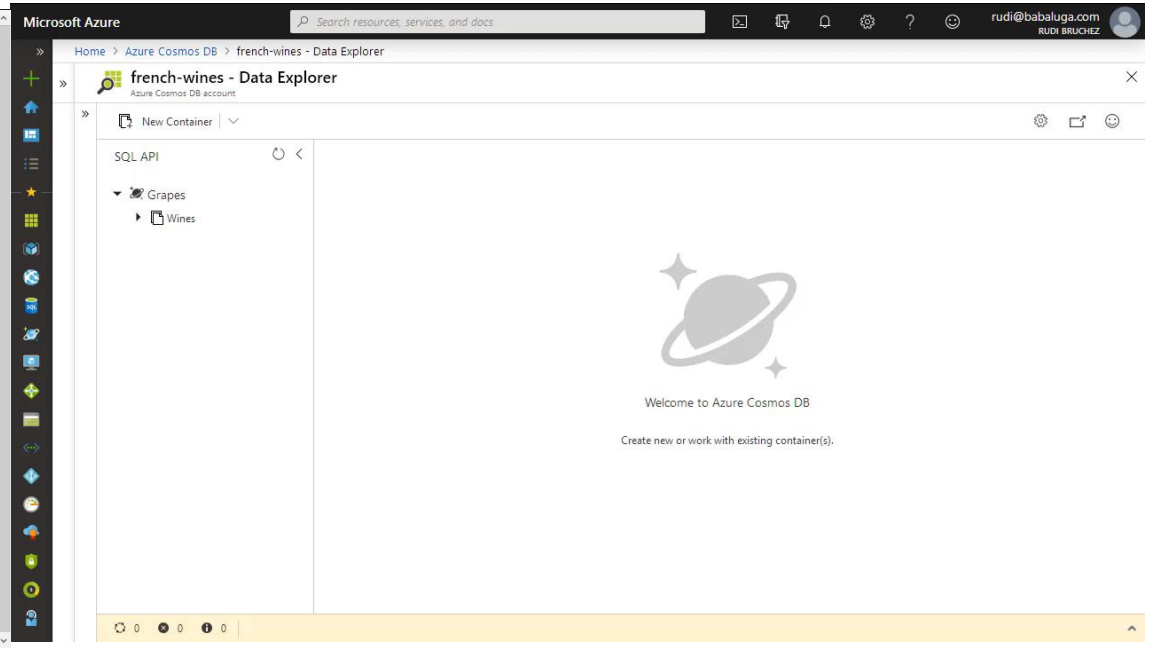

Spatial Data with the SQL API

GeoJSON specification (rfc 7946)

Points, LineStrings, and Polygons

WGS-84 CRS only

- World Geodetic System used by GPS, Google Map, Bing Map

OGC functions : ST_DISTANCE, ST_WITHIN, ST_INTERSECTS

```
{
    "type":"Polygon",
    "coordinates":[ [
        [ 31.8, -5 ],
        [ 31.8, -4.7 ],
        [ 32, -4.7 ],
        [ 32, -5 ],
        [ 31.8, -5 ]
    ] ]
}
```

020.spatial.cs

MongoDB API

JSON Like the SQL API

MongoDB compatible

You can use the MongoDB tools and the mongo clients

Add an _id identifier to be MongoDB compatible

https://docs.microsoft.com/en-us/azure/cosmos-db/mongodb-feature-support

MongoDb Sharding

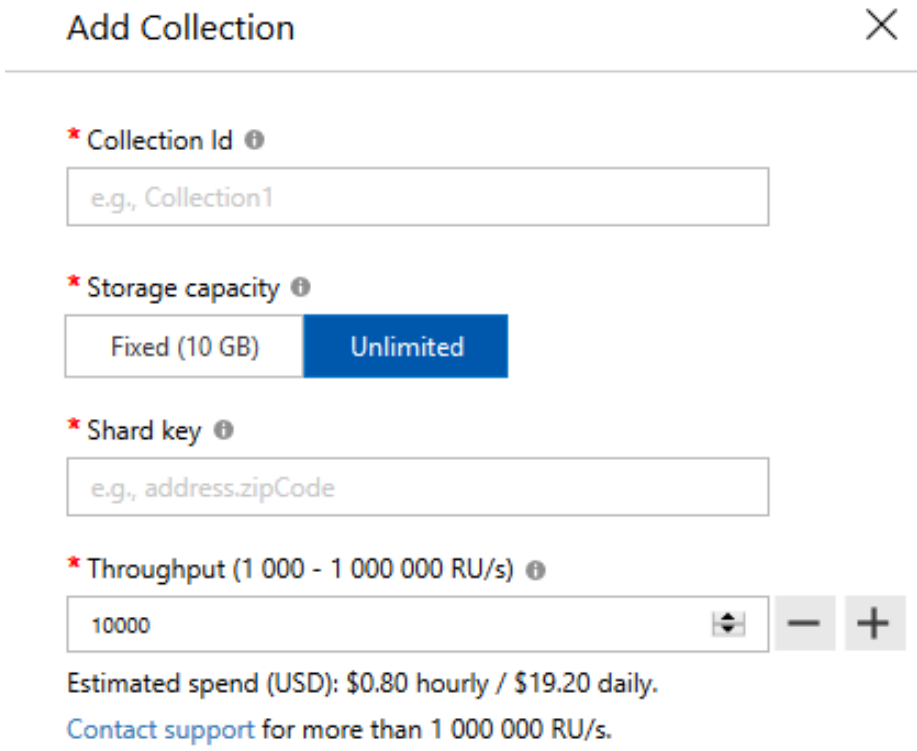

Importing into Mongo CosmosDB

```
./mongoimport.exe --host pachamongo.documents.azure.com:10255 -u pachamongo -p
Sa98P8ahFdcQF5JZI7S3RdAThpkGJCJ8qSSHd51q8JB914ieyler380Q5KQSISb87U1Zmo6k6QfND6e2GM
q6zg== --ssl --sslAllowInvalidCertificates --db pachamongo --collection
restaurants --drop --file ./primer-dataset.json
```

Mongo
MongoImport
MongoExport

Demo – MongoDB API

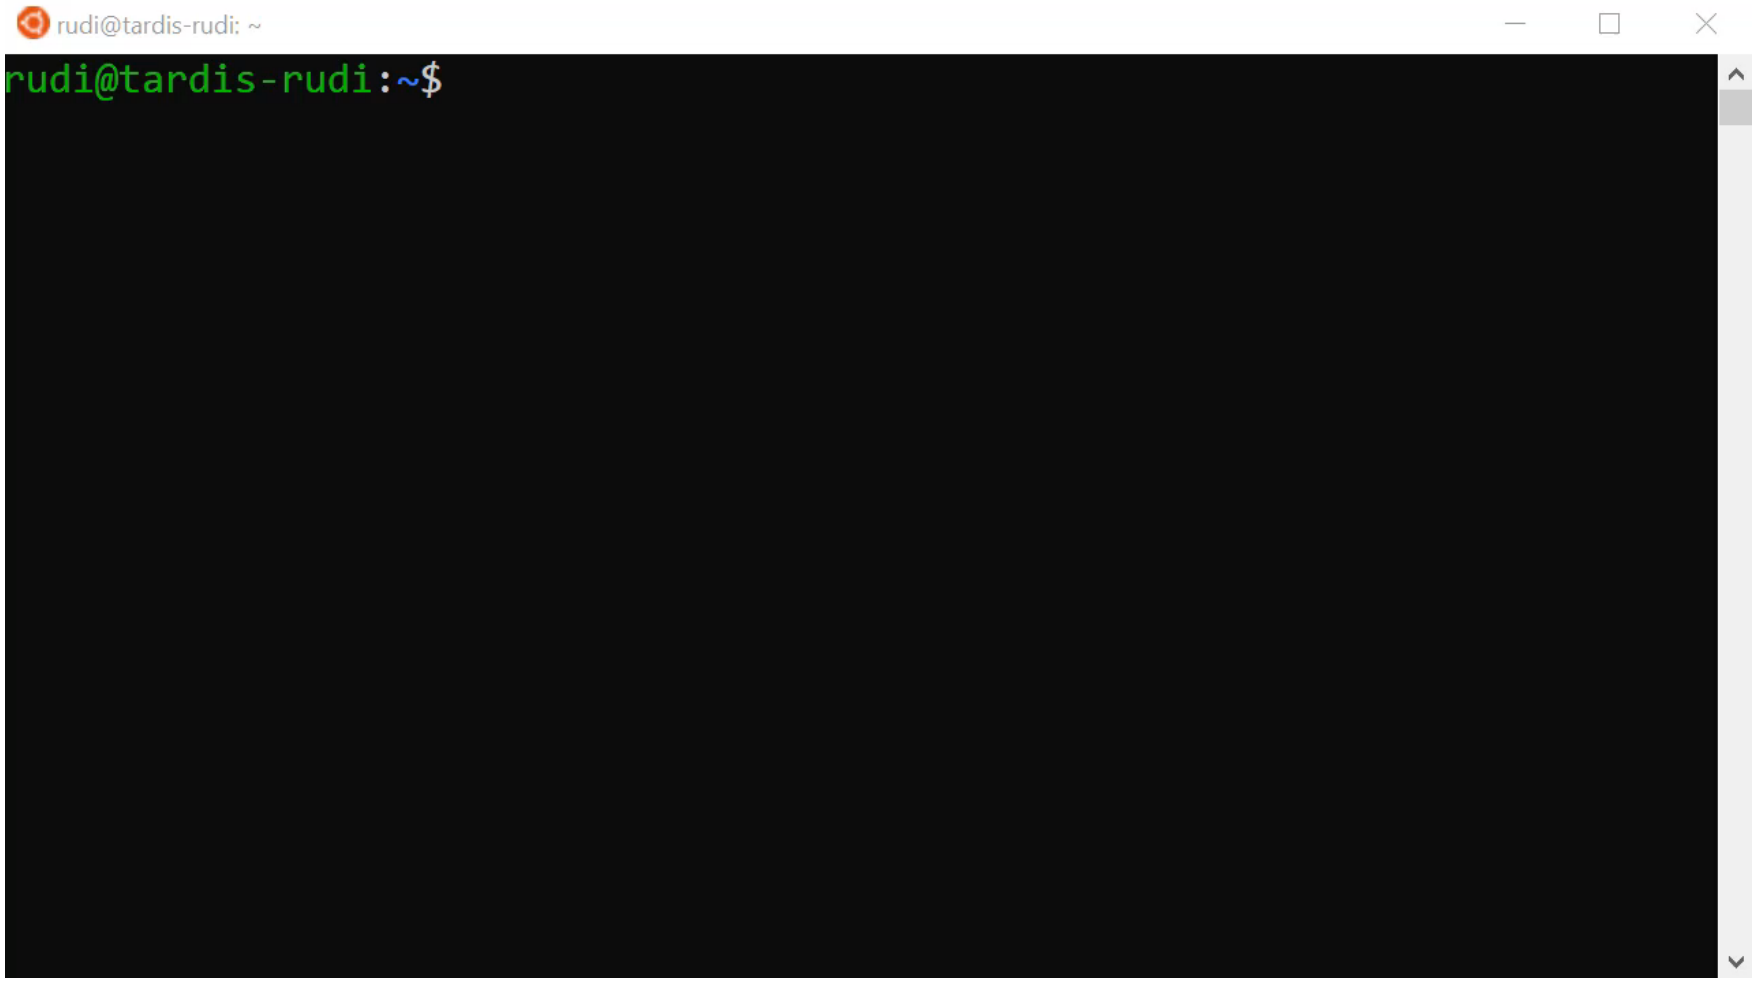

Demo – GUI and Aggregation Framework

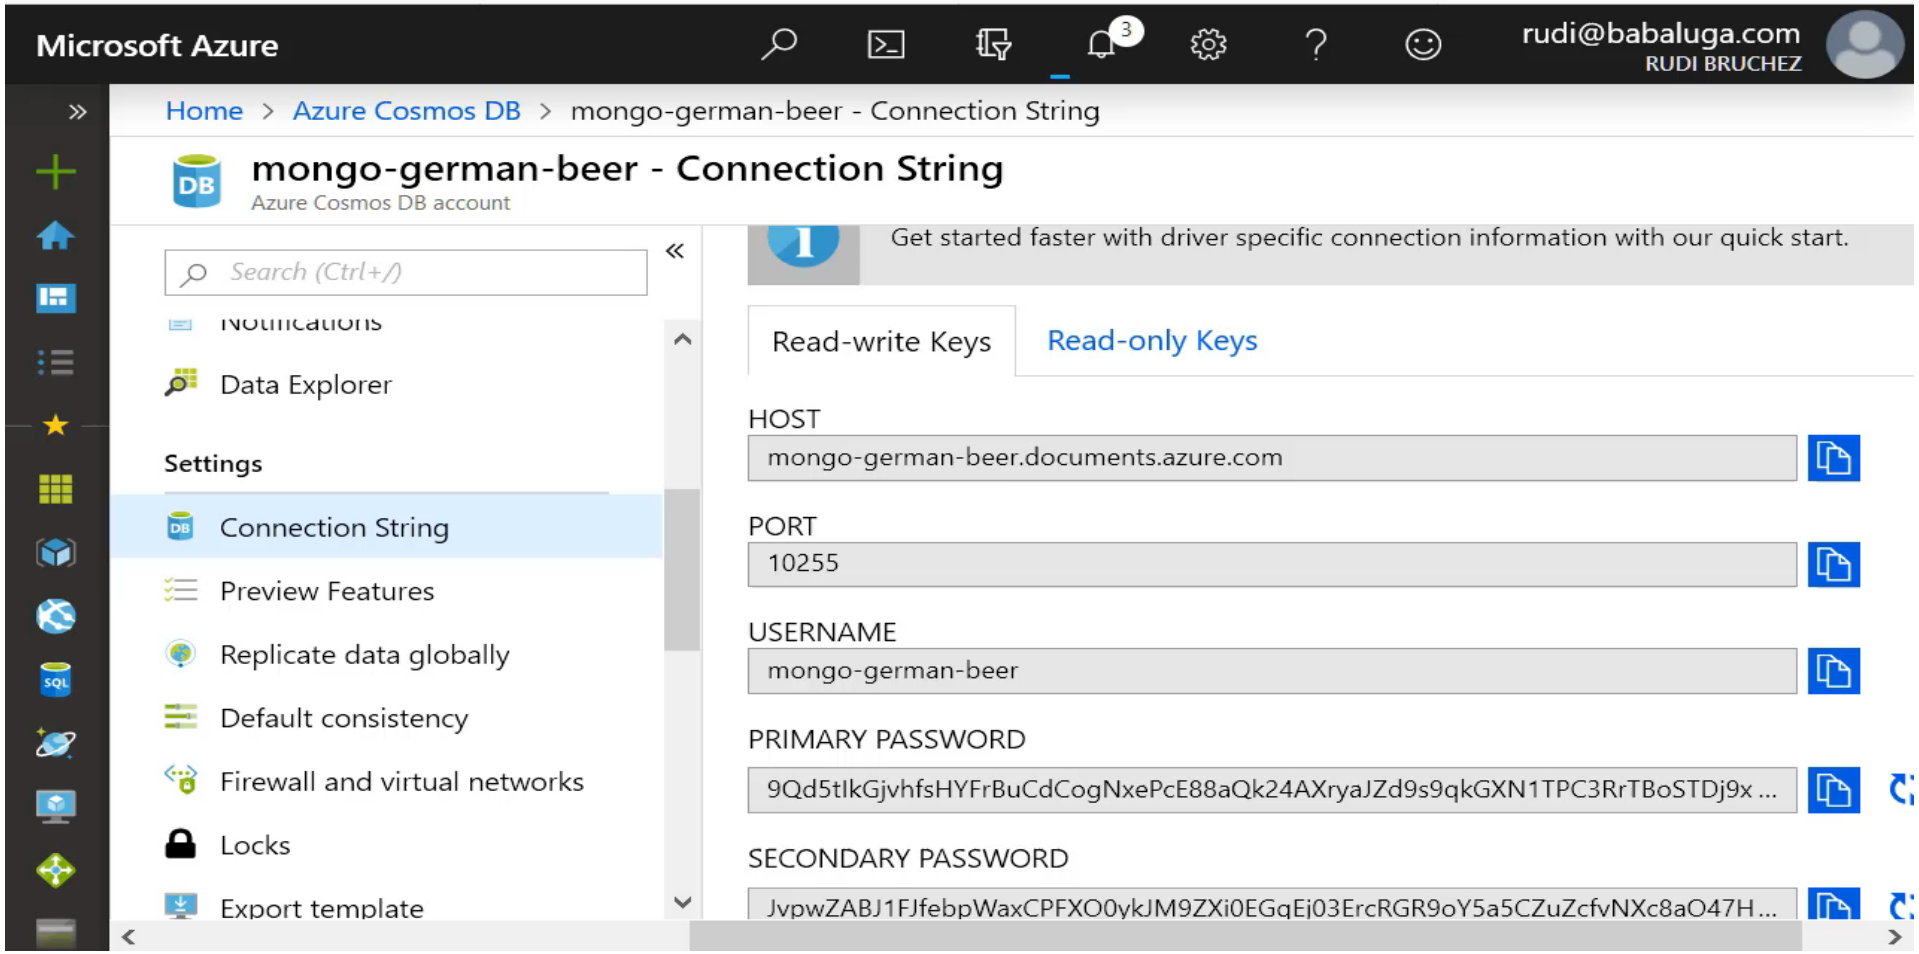

https://community.qlik.com/t5/Technology-Partners-Ecosystem/White-Paper-Connecting-to-CosmosDB-Mongo-API-using-Qlik-MongoDB/ta-p/1527975

Table API

Simple Key-Value (KV) store

Get and Set

You can search in values, everything is indexed

Hash table, very fast for keys

API is recent : https://docs.microsoft.com/en-us/azure/cosmos-db/table-sdk-dotnet

No support yet for .NET Core (use the old one)

Graph API

Based on Tinkerpop, Gremlin language

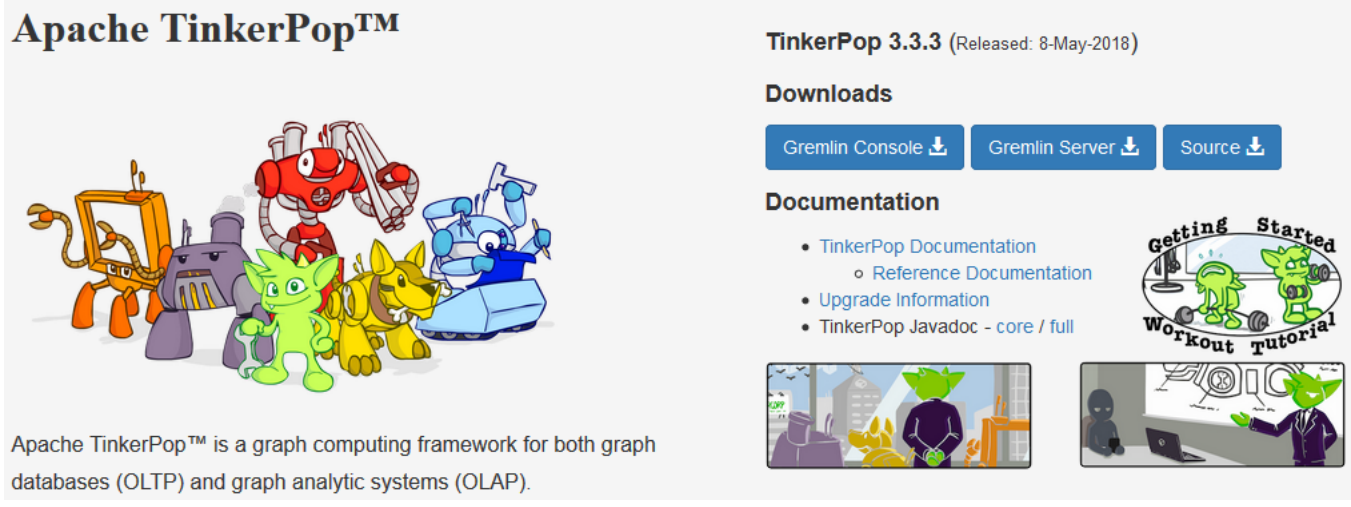

Edge / arc

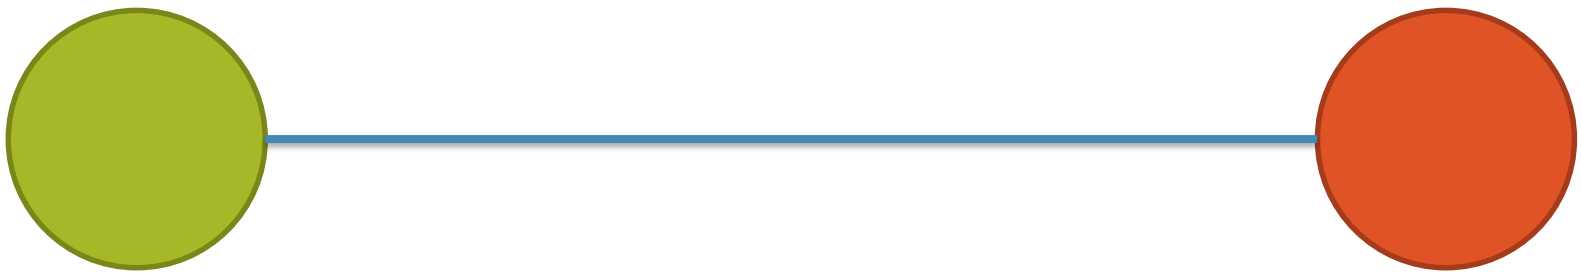

Undirected

Directed

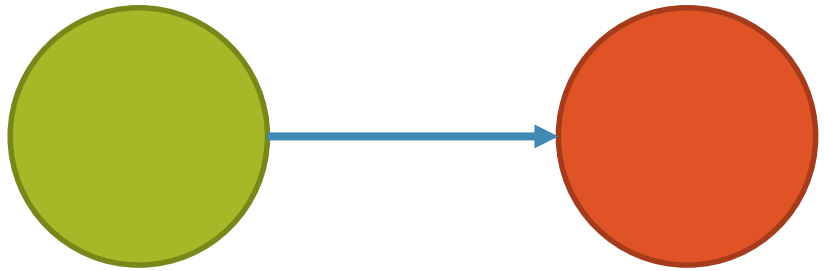

Property graph

A graph database is a property graph

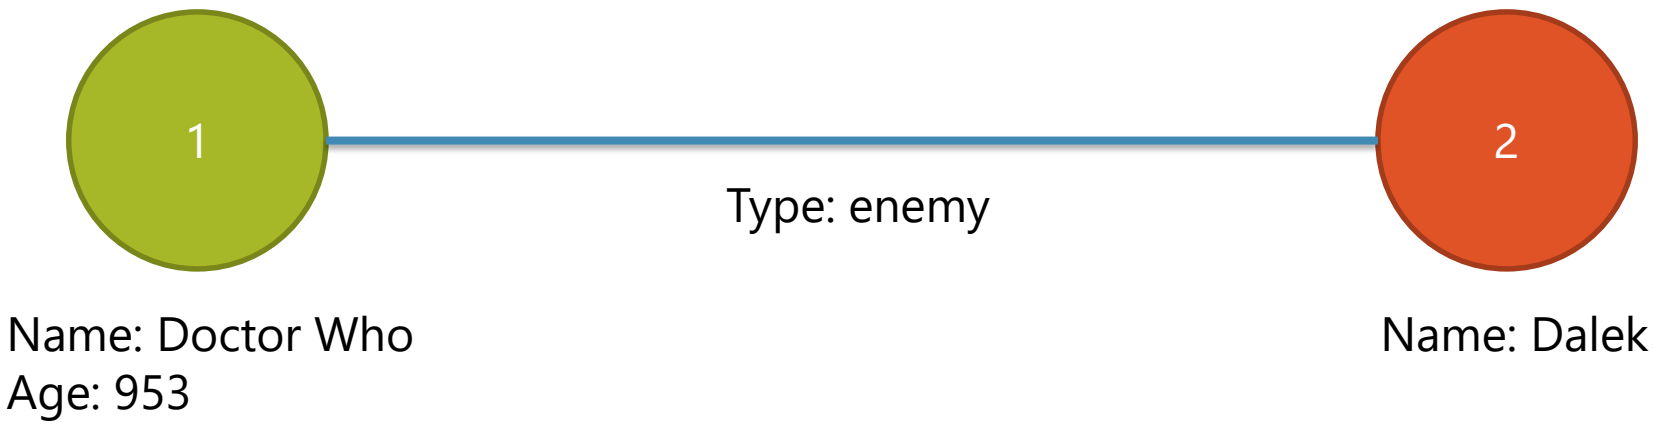

Domain Specific Languages

Cypher

Neo4J

```
MATCH (actor:Person)-[:ACTED_IN]->(movie:Movie)
WHERE movie.title STARTS WITH "T"
RETURN movie.title AS title, collect(actor.name) AS cast
ORDER BY title ASC LIMIT 10;
```

Apache Tinkerpop

```
g.V().as("a").out("knows").as("b").
  select("a","b").
    by("name").
    by("age")
```

040.movie.gremlin

Cassandra API

Apache drivers compliant with CQLv4

All CQL command supported

All Data Types supported

All functions supported

Operations

Documents are stored as a tree

```
{ "locations": [
    { "country": "Germany", "city": "Berlin" },
    { "country": "France", "city": "Paris" }
  ],
  "headquarter": "Belgium",
  "exports":[{ "city": "Moscow" },
             { "city": "Athens"}]
};
```

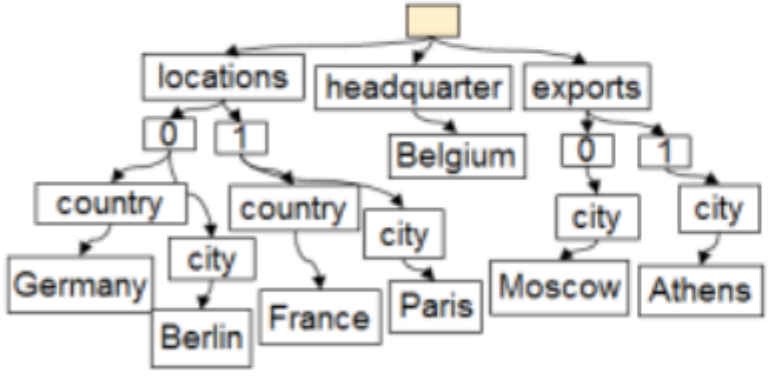

Document 1

```
{ "locations": [
    { "country": "Germany",
      "city": "Bonn","revenue": 200
    }],
  "headquarter": "Italy",
  "exports": [
    { "city": "Berlin","dealers":[{"name": "Hans"
    { "city": "Athens" }]
};
```

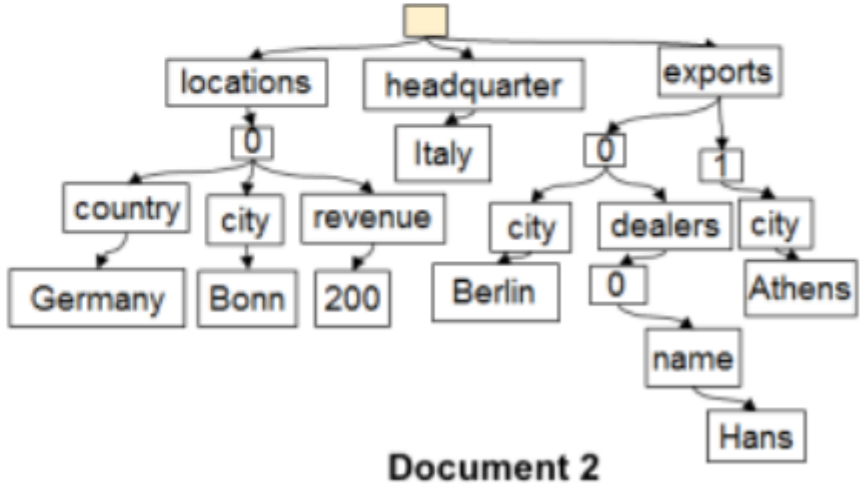

Document 2

Automatic indexing the tree

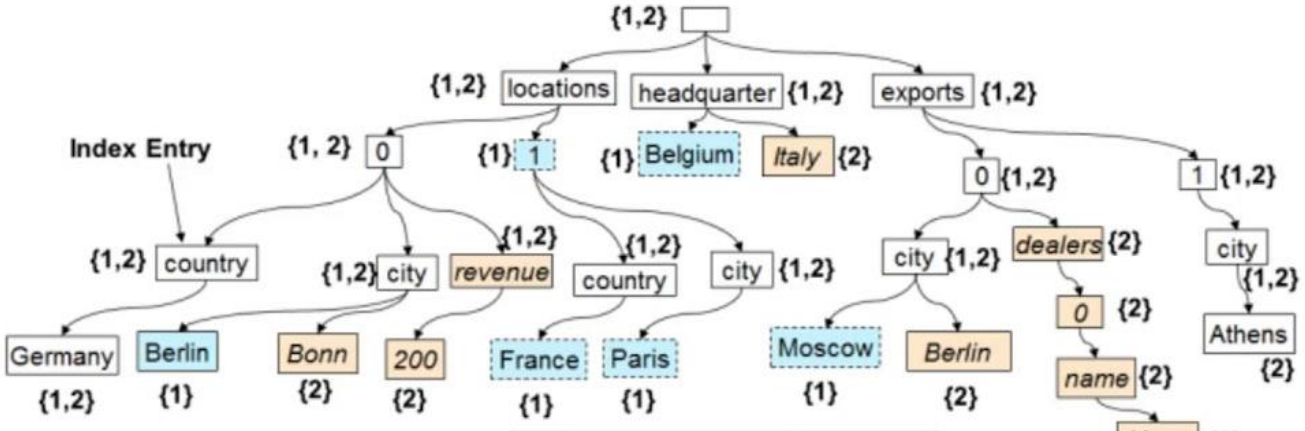

Term	Postings
$/locations/0/	1, 2
locations/0/country/	1, 2
locations/0/city/	1, 2
0/country/Germany	1, 2
1/country/France	1
...	...
0/city/Moscow	1
0/dealers/0	2
...	...

Automatic. All fields are indexed.

Can be set manually, even by document

Online strategy changes, no impact on RUs

Indexing

Consistent: changes happen immediately. higher RU consumption

Lazy: asynchronous changes, background process. Query consistency is eventual and RU consumption is lower.

Index Types

Hash: useful for equality and inequality predicates.

Range: useful for ordering and range searches.

Spatial: useful for spatial queries (within, distance, etc.)

070.Index.cs

Demo – indexing

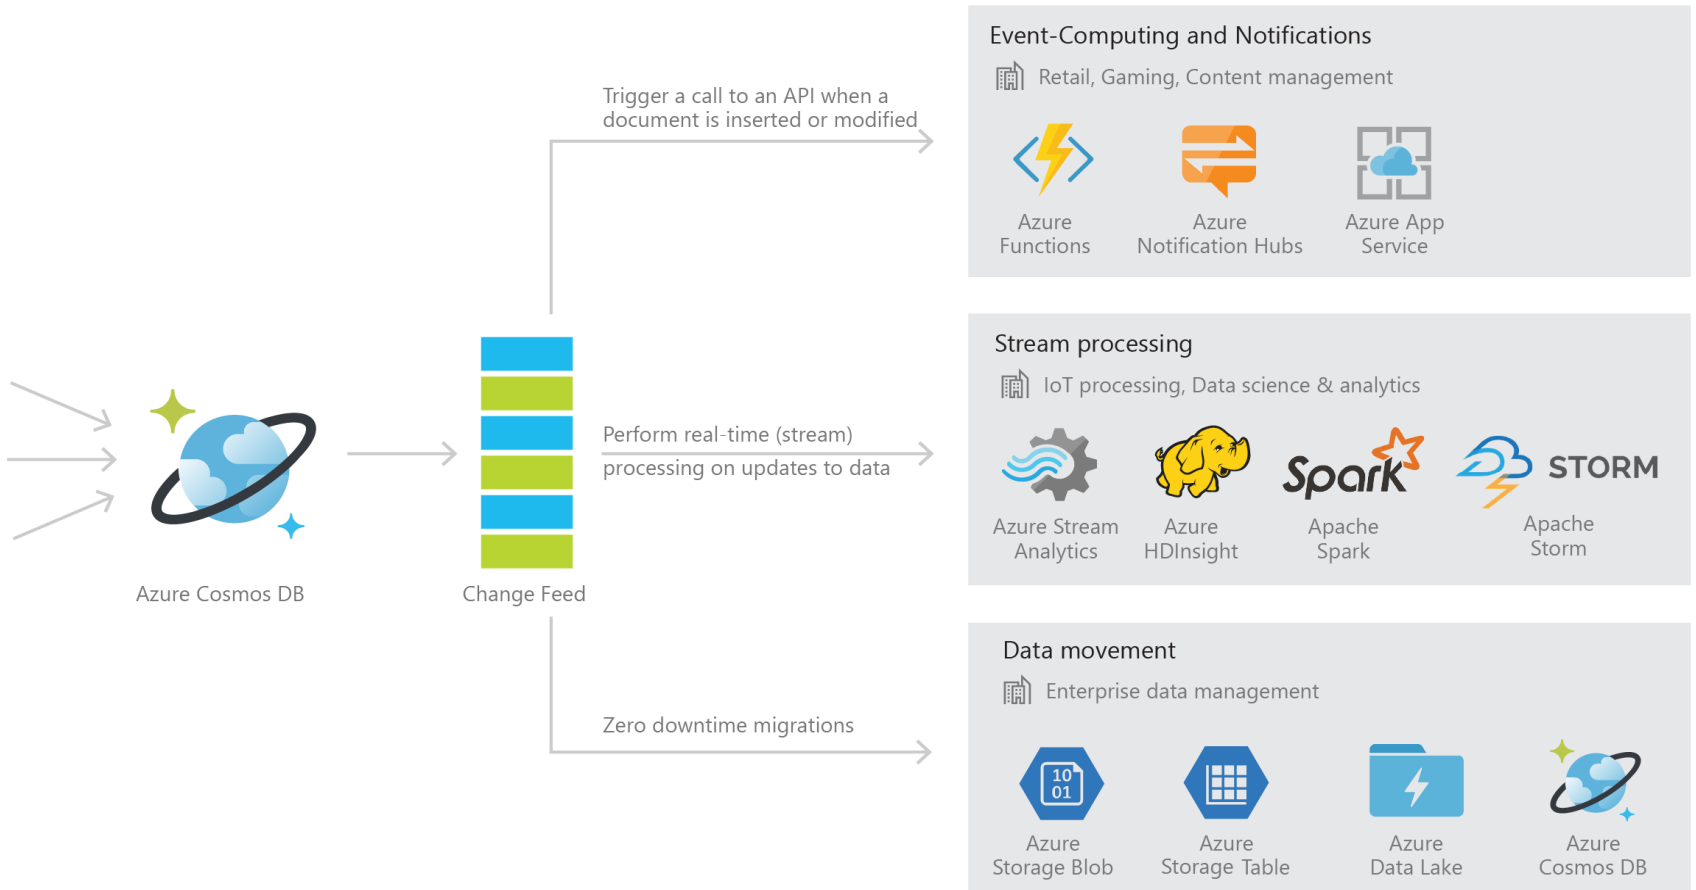

Attachments

For binaries

REST API – POST using AtomPub (rfc 5023)

stored in CosmosDB: POST with the raw attachment as body. 2 headers: Content-Type (MIME type) and Slug (name)

External: post just the attachment metadata

Internally stored: 2 GB limit per account.

Change feed

Supported now for .NET, Java, Python and Node/JS SDKs and for Core and Gremlin APIs.

SP, Triggers, UDF

Javascript

```
UserDefinedFunction regexMatchUdf = new UserDefinedFunction
  {
      Id = "REGEX_MATCH",
      Body = @"function (input, pattern) {
                  return input.match(pattern) !== null;
              };",
  };

  UserDefinedFunction createdUdf = client.CreateUserDefinedFunctionAsync(
      UriFactory.CreateDocumentCollectionUri("testdb", "families"),
      regexMatchUdf).Result;
```

050.ServerSideScripts.cs

Offline emulator

Msi or Docker

Fully supports SQL API and MongoDB collections
Table, Graph, Cassandra not fully supported (yet)

No scalability, obviously

localhost:8081

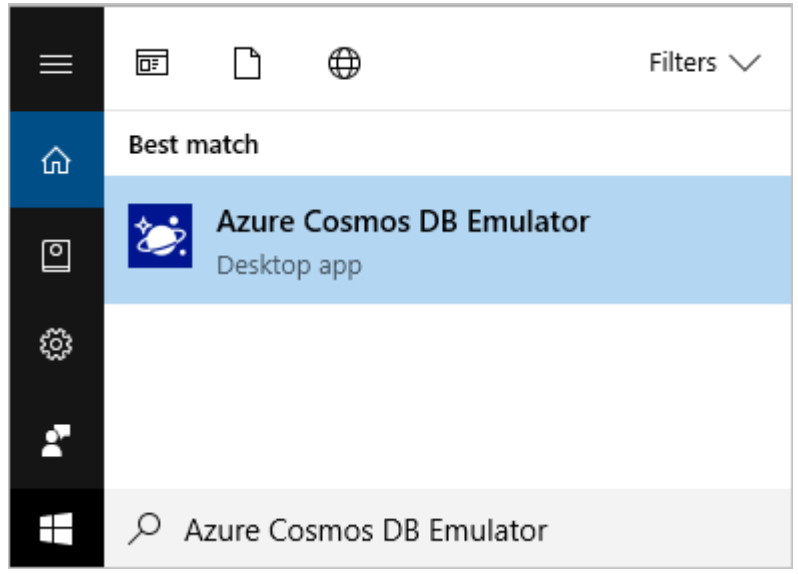

Using the local emulator

Develop for free without an Azure account

Docker :

microsoft/azure-cosmosdb-emulator

Docker Windows Container

Data migration tool

From

JSON files, MongoDB

SQL Server, CSV

Azure Table storage

Amazon DynamoDB

HBase

Azure Cosmos DB collections

To

SQL API – all sources

Table API - Data Migration tool or AzCopy.

MongoDB API - export only, import using MongoDB tools

Graph API - not supported yet

Data migration tool

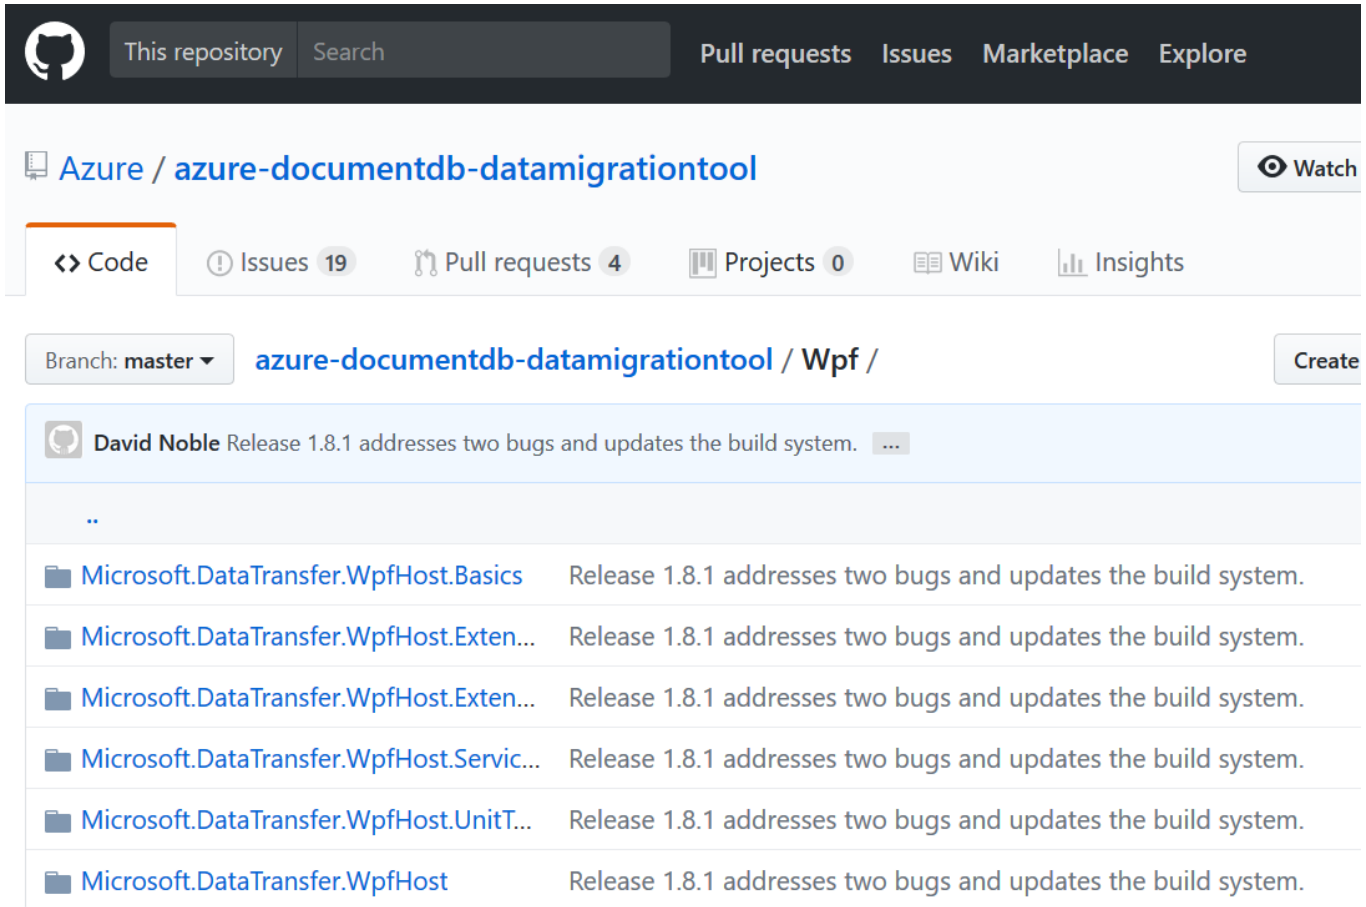

https://github.com/azure/azure-documentdb-datamigrationtool

Change Feed

Event-Computing and Notifications

Retail, Gaming, Content management

- Azure Functions
- Azure Notification Hubs
- Azure App Service

Stream processing

IoT processing, Data science & analytics

- Azure Stream Analytics
- Azure HDInsight
- Apache Spark
- Apache Storm

Data movement

Enterprise data management

- Azure Storage Blob
- Azure Storage Table
- Azure Data Lake
- Azure Cosmos DB

New – Apache Spark Execution

* API

Apache Spark

Cassandra

Enable Disable

Sign up for Apache Spark Preview

Resources

Azure Cosmos DB query cheat sheets

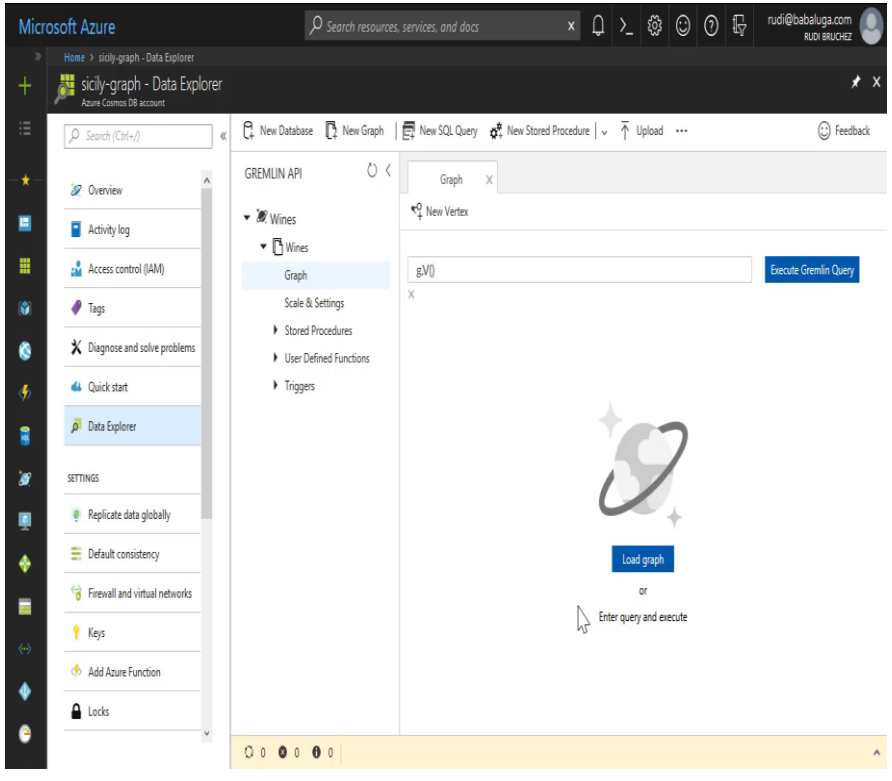

https://docs.microsoft.com/en-us/azure/cosmos-db/query-cheat-sheet

Session Feedback Day
(not optional!)

http://bit.ly/DataGrillen2019Day

Event Feedback (not optional!)

http://bit.ly/DataGrillen2019Ever

That’s all folks!